

```
In [ ]: import tensorflow as tf
        from tensorflow.keras import layers, models

        train_dir = '/Train_Folder_out/'
        validation_dir = '/Validation_Folder_out/'

        img_size = (64, 64)
        batch_size = 32
        train_dataset = tf.keras.preprocessing.image_dataset_from_directory(
            train_dir,
            image_size=img_size,
            batch_size=batch_size,
        )
        validation_dataset = tf.keras.preprocessing.image_dataset_from_directory(
            validation_dir,
            image_size=img_size,
            batch_size=batch_size,
        )
        model = models.Sequential([
            layers.Conv2D(32, (4, 4), activation='relu', input_shape=(64, 64, 3)),
            layers.AveragePooling2D(2, 2),
            layers.Conv2D(64, (3, 3), activation='relu'),
            layers.MaxPooling2D(2, 2),
            layers.Flatten(),
            layers.Dense(128, activation='relu'),
            layers.Dense(3, activation='softmax')
        ])
        model.compile(
            optimizer='adam',
            loss=tf.keras.losses.SparseCategoricalCrossentropy(from_logits=False),
            metrics=['accuracy']
        )
        history = model.fit(
            train_dataset,
            epochs=10,
            validation_data=validation_dataset
        )

        test_loss, test_acc = model.evaluate(validation_dataset)
        print('Test accuracy:', test_acc)
        model.summary()
```

Found 1200 files belonging to 3 classes.

Found 487 files belonging to 3 classes.

Epoch 1/10

38/38 [=====] - 5s 116ms/step - loss: 65.0976 - accuracy: 0.8200 - val_loss: 0.0013 - val_accuracy: 1.0000

Epoch 2/10

38/38 [=====] - 4s 112ms/step - loss: 1.6156e-04 - accuracy: 1.0000 - val_loss: 3.7426e-06 - val_accuracy: 1.0000

Epoch 3/10

38/38 [=====] - 4s 108ms/step - loss: 3.8107e-07 - accuracy: 1.0000 - val_loss: 6.4930e-06 - val_accuracy: 1.0000

Epoch 4/10

38/38 [=====] - 4s 115ms/step - loss: 5.4141e-08 - accuracy: 1.0000 - val_loss: 5.0649e-06 - val_accuracy: 1.0000

Epoch 5/10

38/38 [=====] - 4s 102ms/step - loss: 3.5564e-08 - accuracy: 1.0000 - val_loss: 4.9085e-06 - val_accuracy: 1.0000

Epoch 6/10

38/38 [=====] - 4s 116ms/step - loss: 2.8014e-08 - accuracy: 1.0000 - val_loss: 5.3857e-06 - val_accuracy: 1.0000

Epoch 7/10

38/38 [=====] - 5s 118ms/step - loss: 2.2252e-08 - accuracy: 1.0000 - val_loss: 5.9006e-06 - val_accuracy: 1.0000

Epoch 8/10

38/38 [=====] - 5s 124ms/step - loss: 1.7881e-08 - accuracy: 1.0000 - val_loss: 6.0246e-06 - val_accuracy: 1.0000

Epoch 9/10

38/38 [=====] - 5s 116ms/step - loss: 1.4901e-08 - accuracy: 1.0000 - val_loss: 6.8982e-06 - val_accuracy: 1.0000

Epoch 10/10

38/38 [=====] - 5s 130ms/step - loss: 1.2418e-08 - accuracy: 1.0000 - val_loss: 7.4136e-06 - val_accuracy: 1.0000

16/16 [=====] - 1s 28ms/step - loss: 7.4136e-06 - accuracy: 1.0000

Test accuracy: 1.0

Model: "sequential_6"

| Layer (type) | Output Shape | Param # |
|---|--------------------|---------|
| ===== | | |
| conv2d_18 (Conv2D) | (None, 61, 61, 32) | 1568 |
| average_pooling2d_12 (AveragePooling2D) | (None, 30, 30, 32) | 0 |
| conv2d_19 (Conv2D) | (None, 28, 28, 64) | 18496 |
| max_pooling2d_6 (MaxPooling2D) | (None, 14, 14, 64) | 0 |
| flatten_6 (Flatten) | (None, 12544) | 0 |
| dense_18 (Dense) | (None, 128) | 1605760 |
| dense_19 (Dense) | (None, 3) | 387 |
| ===== | | |

Total params: 1,626,211
 Trainable params: 1,626,211
 Non-trainable params: 0

```
In [ ]: import tensorflow as tf
from tensorflow.keras import layers, models

train_dir = '/Train_Folder_out/'
validation_dir = '/Test_Folder_out/'

img_size = (64, 64)
batch_size = 32

train_dataset = tf.keras.preprocessing.image_dataset_from_directory(
    train_dir,
    image_size=img_size,
    batch_size=batch_size,
)
validation_dataset = tf.keras.preprocessing.image_dataset_from_directory(
    validation_dir,
    image_size=img_size,
    batch_size=batch_size,
)
model = models.Sequential([
    layers.Conv2D(32, (3, 3), activation='relu', input_shape=(64, 64, 3)),
    layers.AveragePooling2D(2, 2),
    layers.Conv2D(64, (3, 3), activation='relu'),
    layers.AveragePooling2D(2, 2),
    layers.Conv2D(128, (3, 3), activation='relu'),
    layers.MaxPooling2D(2, 2),
    layers.Flatten(),
    layers.Dense(128, activation='relu'),
    layers.Dense(64, activation='relu'),
    layers.Dense(3, activation='softmax')
])
model.compile(
    optimizer='adam',
    loss=tf.keras.losses.SparseCategoricalCrossentropy(from_logits=False),
    metrics=['accuracy']
)
early_stopping = tf.keras.callbacks.EarlyStopping(
    monitor='val_loss',
    patience=3,
    restore_best_weights=True
)
model_checkpoint = tf.keras.callbacks.ModelCheckpoint(
    "best_model.h5",
    monitor='val_loss',
    save_best_only=True
)
history = model.fit(
    train_dataset,
    epochs=10,
    validation_data=validation_dataset,
```

```
        callbacks=[early_stopping, model_checkpoint]
    )
    best_model = tf.keras.models.load_model("best_model.h5")
    test_loss, test_acc = best_model.evaluate(validation_dataset)
    print('Test accuracy:', test_acc)
    best_model.summary()
```

Found 1200 files belonging to 3 classes.

Found 487 files belonging to 3 classes.

Epoch 1/10

38/38 [=====] - 6s 122ms/step - loss: 3.8176 - accuracy: 0.8808 - val_loss: 0.0013 - val_accuracy: 1.0000

Epoch 2/10

38/38 [=====] - 5s 121ms/step - loss: 7.3711e-05 - accuracy: 1.0000 - val_loss: 9.4498e-05 - val_accuracy: 1.0000

Epoch 3/10

38/38 [=====] - 5s 120ms/step - loss: 1.4453e-05 - accuracy: 1.0000 - val_loss: 8.3171e-06 - val_accuracy: 1.0000

Epoch 4/10

38/38 [=====] - 5s 118ms/step - loss: 6.9727e-07 - accuracy: 1.0000 - val_loss: 1.1828e-05 - val_accuracy: 1.0000

Epoch 5/10

38/38 [=====] - 5s 118ms/step - loss: 4.9988e-07 - accuracy: 1.0000 - val_loss: 1.4547e-05 - val_accuracy: 1.0000

Epoch 6/10

38/38 [=====] - 5s 120ms/step - loss: 4.1465e-07 - accuracy: 1.0000 - val_loss: 1.7004e-05 - val_accuracy: 1.0000

16/16 [=====] - 1s 30ms/step - loss: 8.3171e-06 - accuracy: 1.0000

Test accuracy: 1.0

Model: "sequential_7"

| Layer (type) | Output Shape | Param # |
|---|---------------------|---------|
| ===== | | |
| conv2d_20 (Conv2D) | (None, 62, 62, 32) | 896 |
| average_pooling2d_13 (AveragePooling2D) | (None, 31, 31, 32) | 0 |
| conv2d_21 (Conv2D) | (None, 29, 29, 64) | 18496 |
| average_pooling2d_14 (AveragePooling2D) | (None, 14, 14, 64) | 0 |
| conv2d_22 (Conv2D) | (None, 12, 12, 128) | 73856 |
| max_pooling2d_7 (MaxPooling2D) | (None, 6, 6, 128) | 0 |
| flatten_7 (Flatten) | (None, 4608) | 0 |
| dense_20 (Dense) | (None, 128) | 589952 |
| dense_21 (Dense) | (None, 64) | 8256 |
| dense_22 (Dense) | (None, 3) | 195 |

=====

Total params: 691,651

Trainable params: 691,651

Non-trainable params: 0

```
In [ ]: import tensorflow as tf
        from tensorflow import keras
        from sklearn.metrics import classification_report

        best_model = tf.keras.models.load_model("best_model.h5")
        validation_loss, validation_accuracy = best_model.evaluate(validation_dataset, verbose=0)
        print('Validation accuracy:', validation_accuracy)

        # Save the best model to a custom location
        save_path = "/examples/"
        best_model.save(save_path)

        true_labels = []
        predicted_labels = []
        for images, labels in validation_dataset:
            true_labels.extend(labels.numpy())
            predicted_labels.extend(tf.argmax(best_model.predict(images), axis=-1).numpy())
        report = classification_report(true_labels, predicted_labels)

        print("Classification Report:")
        print(report)
```

16/16 - 1s - loss: 8.3171e-06 - accuracy: 1.0000 - 917ms/epoch - 57ms/step
Validation accuracy: 1.0

WARNING:absl:Found untraced functions such as _jit_compiled_convolution_op, _jit_compiled_convolution_op, _jit_compiled_convolution_op, _update_step_xla while saving (showing 4 of 4). These functions will not be directly callable after loading.

INFO:tensorflow:Assets written to: /examples/assets

INFO:tensorflow:Assets written to: /examples/assets

```

1/1 [=====] - 0s 86ms/step
1/1 [=====] - 0s 86ms/step
1/1 [=====] - 0s 51ms/step
1/1 [=====] - 0s 51ms/step
1/1 [=====] - 0s 48ms/step
1/1 [=====] - 0s 51ms/step
1/1 [=====] - 0s 50ms/step
1/1 [=====] - 0s 51ms/step
1/1 [=====] - 0s 51ms/step
1/1 [=====] - 0s 51ms/step
1/1 [=====] - 0s 51ms/step
1/1 [=====] - 0s 52ms/step
1/1 [=====] - 0s 52ms/step
1/1 [=====] - 0s 51ms/step
1/1 [=====] - 0s 50ms/step
1/1 [=====] - 0s 70ms/step
1/1 [=====] - 0s 72ms/step

```

Classification Report:

| | precision | recall | f1-score | support |
|--------------|-----------|--------|----------|---------|
| 0 | 1.00 | 1.00 | 1.00 | 164 |
| 1 | 1.00 | 1.00 | 1.00 | 166 |
| 2 | 1.00 | 1.00 | 1.00 | 157 |
| accuracy | | | 1.00 | 487 |
| macro avg | 1.00 | 1.00 | 1.00 | 487 |
| weighted avg | 1.00 | 1.00 | 1.00 | 487 |

```

In [ ]: import matplotlib.pyplot as plt
import numpy as np
num_samples = 15
test_images = []
test_labels = []
for images, labels in validation_dataset.take(num_samples):
    test_images.append(images[0])
    test_labels.append(labels[0].numpy())

if len(test_images) < num_samples:
    num_samples = len(test_images)
    print(f"Reduced num_samples to {num_samples} because the dataset contains fewer


predicted_labels = best_model.predict(np.array(test_images))
predicted_labels = tf.argmax(predicted_labels, axis=-1).numpy()

class_names = train_dataset.class_names
if num_samples > 0:
    plt.figure(figsize=(15, 15))
    for i in range(num_samples):
        plt.subplot(5, 5, i+1)
        plt.imshow(test_images[i].numpy().astype("uint8"))
        true_label = class_names[test_labels[i]]
        predicted_label = class_names[predicted_labels[i]]
        plt.title(f"True Label: {true_label}\nPredicted Label: {predicted_label}")
        plt.axis('off')
    plt.tight_layout()

```

```
plt.show()
else:
    print("No samples to plot.")
```

1/1 [=====] - 0s 35ms/step



```
In [ ]: import tensorflow as tf
from tensorflow import keras
import matplotlib.pyplot as plt
best_model = tf.keras.models.load_model("best_model.h5")
validation_loss, validation_accuracy = best_model.evaluate(validation_dataset)
print('Validation accuracy:', validation_accuracy)
plt.figure(figsize=(10, 5))
plt.plot(history.history['loss'], label='Training Loss')
plt.plot(history.history['val_loss'], label='Validation Loss')
plt.title('Training and Validation Loss')
plt.xlabel('Epoch')
plt.ylabel('Loss')
plt.legend()
plt.show()
plt.figure(figsize=(10, 5))
plt.plot(history.history['accuracy'], label='Training Accuracy')
plt.plot(history.history['val_accuracy'], label='Validation Accuracy')
plt.title('Training and Validation Accuracy')
plt.xlabel('Epoch')
plt.ylabel('Accuracy')
plt.legend()
plt.show()
```

16/16 [=====] - 1s 32ms/step - loss: 8.3171e-06 - accuracy: 1.0000
Validation accuracy: 1.0

